Bayesian Neural Network via Stochastic Gradient Descent

Abhinav Sagar

Vellore Institute of Technology, Vellore, Tamil Nadu, India

ABSTRACT

The goal of bayesian approach used in variational inference is to minimize the KL divergence between variational distribution and unknown posterior distribution. This is done by maximizing the Evidence Lower Bound (ELBO). A neural network is used to parametrize these distributions using Stochastic Gradient Descent. This work extends the work done by others by deriving the variational inference models. We show how SGD can be applied on bayesian neural networks by gradient estimation techniques. For validation, we have tested our model on 5 UCI datasets and the metrics chosen for evaluation are Root Mean Square Error (RMSE) error and negative log likelihood. Our work considerably beats the previous state of the art approaches for regression using bayesian neural networks.

1. Introduction

Recently, there has been a lot of work done on inference using probabilistic models. In this approach, rather than considering the parameters of the neural network as point estimates, we sample them as continuous distributions. Using this approach, helps us infer the uncertainty involved while making the predictions. This is very important in sensitive domains where not only we want to find out the predictions made by the model but also with how much certainty it is making the predictions.

The problem with this approach lies in the calculation of posterior distribution which is often intractable. Hence for the computation, it is necessary to convert the variational distribution into a tractable posterior distribution. Variational inference approach is used to convert the inference problem into an optimization problem with the objective of minimizing the KL-divergence between variational distribution and true posterior. This is done by maximizing the ELBO.

In this paper, we present the background of a Bayesian Neural Network using stochastic gradient descent. Then we show how distributions can be parameterized by using variational inference techniques. We validated our work on UCI datasets and show our approach is better than the previous results obtained.

2. Background

2.1 Variational Inference and the ELBO

A probabilistic model is denoted using observations x, latent variables z and model parameters θ. The optimal θ value has to be found to maximize the marginal likelihood as given in below equation where we refer to $p_\theta(x \mid z)$ as the generative distribution and to $p_\theta(z)$ as the prior distribution.

$$p_\theta(\mathbf{x}) = \int p_\theta(\mathbf{x}|\mathbf{z})p_\theta(\mathbf{z})\mathrm{d}\mathbf{z}$$

Computing the posterior by doing inference is intractable as it requires doing an integration over z. Hence we maximize Evidence Lower Bound (ELBO) in variational inference which can be computed by approximating posterior on the latent variable as shown in the below equation.

$$\begin{aligned}
\ln p_\theta(\mathbf{x}) &= \ln \int p_\theta(\mathbf{x}|\mathbf{z})p_\theta(\mathbf{z})\mathrm{d}\mathbf{z} \\
&= \ln \int \frac{q_\phi(\mathbf{z}|\mathbf{x})}{q_\phi(\mathbf{z}|\mathbf{x})} p_\theta(\mathbf{x}|\mathbf{z})p_\theta(\mathbf{z})\mathrm{d}\mathbf{z} \\
&= \ln \int q_\phi(\mathbf{z}|\mathbf{x}) \frac{p_\theta(\mathbf{x}|\mathbf{z})p_\theta(\mathbf{z})}{q_\phi(\mathbf{z}|\mathbf{x})}\mathrm{d}\mathbf{z}
\end{aligned}$$

2.2 Stochastic Gradient Descent

We can optimize the variational parameters φ by using stochastic gradient ascent with the following update rule as given in the below equation. Here γ is learning rate and n is the size of randomly sampled mini batches of training data.

$$\phi_{t+1} = \phi_t + \gamma \sum_{i=1}^{n} \nabla_\phi \mathcal{L}\left(\mathbf{x}^{(i)}\right)$$

We define ESBO S(λ) having its posterior unknown. Hence we use joint distribution instead. The weights are normalized to remove the intractable integration present in the denominator. The dataset has been divided into mini-batches and the operations are carried on them in successive iterations. The ESBO is defined in the below equation.

$$\hat{\nabla}_\lambda \mathcal{S} = \sum_{i}^{M} \hat{w}^{(i)} \left[\left(\log \omega\left(\theta^{(i)}\right) + 1\right) \nabla_\lambda \log p\left(\mathcal{D}, \theta^{(i)}\right) - \nabla_\lambda \log q\left(\theta^{(i)}; \lambda\right)\right]$$

Using the reparameterization trick, the gradient can be written as shown in Fig 4. This operation is very helpful in reducing the complexity of variational inference models as shown in the below equation.

$$\hat{\nabla}_\lambda \mathcal{S} = \sum_{i}^{M} \hat{w}^{(i)} \left[\left(\log \omega\left(\theta^{(i)}\right) + 1\right) \nabla_\lambda \log p\left(\mathcal{D}, g_\lambda\left(\epsilon^{(i)}\right)\right) - \nabla_\lambda \log q\left(g_\lambda\left(\epsilon^{(i)}\right); \lambda\right)\right]$$

Now Stochastic Gradient Descent (SGD) can be applied to minimize the ESBO. This operation is shown in the equation below.

$$\hat{\mathcal{S}}^* = \sum_{i=1}^{M} \hat{w}^{(i)} \left[\sum_{n=1}^{N} \log p\left(x_n|\theta^{(i)}\right) + \log p\left(\theta^{(i)}\right) - \log q\left(\theta^{(i)}; \lambda^*\right)\right]$$

2.3 Algorithm

Next we present the algorithm used in this work:

λ_0 and the learning rate α_0 are initialized using arbitrary values

for t=0, ..., T

    generate M samples $\left\{\theta^{(i)}\right\}_{i=1}^{M} : \epsilon^{(i)} \sim \mathcal{N}(0,1), \theta^{(i)} = g_\lambda \left(\epsilon^{(i)} = \mu + \sigma\epsilon^{(i)}\right)$

    calculate the weight $\log w^{(i)} = \dfrac{N}{S} \sum_{n=1}^{S} \log p\left(x_n|\theta^{(i)}\right) + \log p\left(\theta^{(i)}\right) - \log q\left(\theta^{(i)}; \lambda_t\right)$

    evaluate the gradient $\hat{\nabla}_\lambda \mathcal{S} = -\sum \hat{w}^{(i)} \nabla_\lambda \log q\left(g_\lambda\left(\epsilon^{(i)}\right)\right)$

    $w^{(i)} = \exp\left(\log w^{(i)} + \min\left\{\log w^{(i)}\right\}\right)$

    update $\lambda_{t+1} = \lambda_t - \alpha_t * \nabla_\lambda \mathcal{S}$

3. Simulation Studies

For bayesian neural network regression, we have used datasets from UCI repository: Boston, Concrete, Energy, Protein, Wine. We have used a neural network with one hidden layer with 50 neurons in each case. We set $\eta\,(0, 1)$ as the prior distribution for the weight and bias of the neural network, ReLu as the activation function and batch size as 32. The datasets are randomly partitioned into 90% with training data and 10% for testing, and the results are averaged over 50 random trials. The average RMSE loss and average log likelihood values are given in Table 1 and Table 2 respectively. Our method archives much better results on both the above metrics compared to the previous state of the art.

Table 1: Bayesian neural network regression: average test RMSE(lower is better)

	Ours	Rényi-VI[21]	CLBO-VI [21]	ELBO-VI[15]	BPB[6]
Boston	**2.58±0.13**	2.86±0.40	2.71±0.29	2.89±0.17	2.977±0.093
Concrete	**4.79±0.36**	5.15±0.25	5.04±0.27	5.42±0.11	5.506±0.103
Energy	0.74±0.08	1.00±0.18	0.95±0.15	**0.51±0.01**	1.734±0.051
Protein	**4.38±0.07**	4.65±0.07	4.43±0.05	4.45±0.02	4.623±0.009
Wine	**0.59±0.04**	0.62±0.03	0.61±0.03	0.63±0.01	0.614±0.008

Table 2: Bayesian neural network regression: average negative test LL(lower is better)

	Ours	Rényi-VI[13]	CLBO-VI [13]	ELBO-VI[12]	BPB[11]
Boston	**2.36±0.17**	2.46±0.16	2.40±0.09	2.52±0.03	2.579±0.052
Concrete	**2.93±0.07**	3.04±0.07	3.02±0.03	3.11±0.02	3.137±0.021
Energy	1.53±0.04	1.67±0.05	1.65±0.04	**0.77±0.02**	1.981±0.023
Protein	**2.85±0.02**	2.93±0.00	2.89±0.01	2.91±0.00	2.950±0.002
Wine	**0.92±0.03**	0.94±0.04	0.93±0.04	0.96±0.01	0.931±0.014

4. Conclusion

In this work, we showed how a Bayesian neural network can be trained using stochastic gradient descent. We started with presenting the problem in variational inference approaches and how to convert the problem into a tractable one using ELBO. Next we presented our algorithm which used gradient estimation techniques for doing the inference. We evaluated our work on UCI datasets and this method produces better results than the previous state of the art.

References

1. Adji B. Dieng, Dustin Tran, Rajesh Ranganath, John Paisley, and David M. Blei. Variational inference via χ-upper bound minimization. In Proceedings of the Neural Information Processing Systems, 2017.

2. Justin Domke and Daniel Sheldon. Importance weighting and variational inference. In Proceedings of the Neural Information Processing Systems, 2018.

3. José Miguel Hernández-Lobato and Ryan P. Adams. Probabilistic backpropagation for scalable learning of Bayesian neural networks. In Proceedings of the International Conference on Machine Learning, 2015.

4. Matthew D. Hoffman, David M. Blei, Chong Wang, and John William Paisley. Stochastic variational inference. Journal of Machine Learning Research.

5. Cheng Zhang, Judith Butepage, Hedvig Kjellstrom, and Stephan Mandt. Advances in variational inference. ArXiv preprint, arXiv:1711.05597, 2018.

6. Chenyang Tao, Liqun Chen, Ruiyi Zhang, Ricardo Henao, and Lawrence Carin. Variational inference and model selection with generalized evidence bounds. In Proceedings of the International Conference on Machine Learning, 2018.

7. K. P. Murphy. Machine Learning: A Probabilistic Perspective. MIT press, 2012.

8. Y. Li and R. E. Turner. Variational inference with Rényi divergence. In Proceedings of the Neural Information Processing Systems, 2016.

9. D. Kingma and M. Welling. Auto-encoding variational Bayes. In Proceedings of the 2nd International Conference on Learning Representations (ICLR), 2014.

10. Diederik Kingma and Jimmy Ba. Adam: A method for stochastic optimization. arXiv preprint, arXiv:1412.698, 2014.

11. José Miguel Hernández-Lobato and Ryan P. Adams. Probabilistic backpropagation for scalable learning of Bayesian neural networks. In Proceedings of the International Conference on Machine Learning, 2015.

12. Y. Li and R. E. Turner. Variational inference with Rényi divergence. In Proceedings of the Neural Information Processing Systems, 2016.

13. Chenyang Tao, Liqun Chen, Ruiyi Zhang, Ricardo Henao, and Lawrence Carin. Variational inference and model selection with generalized evidence bounds. In Proceedings of the International Conference on Machine Learning, 2018.