
Stochastic Bayesian Neural Networks

Abhinav Sagar*

Vellore Institute of Technology

Vellore, Tamil Nadu, India

abhinavsagar4@gmail.com

Abstract

Bayesian neural networks perform variational inference over the weights however calculation of the posterior distribution remains a challenge. Our work builds on variational inference techniques for bayesian neural networks using the original Evidence Lower Bound. In this paper, we present a stochastic bayesian neural network in which we maximize Evidence Lower Bound using a new objective function which we name as Stochastic Evidence Lower Bound. We evaluate our network on 5 publicly available UCI datasets using test RMSE and log likelihood as the evaluation metrics. We demonstrate that our work not only beats the previous state of the art algorithms but is also scalable to larger datasets.

1 Introduction

Neural Networks have been highly successful in a variety of domains including computer vision, natural language processing, recommendation systems, reinforcement learning etc. They have considerably surpassed previous state of the art algorithms in machine learning which requires manual feature engineering. However, applying them to sensitive domains like self driving cars, healthcare etc is still a major challenge. This is due to the fact that not only we need predictions made by the model but also with how much certainty it is making those predictions. This is why bayesian neural networks have gained a huge traction recently as they combine the flexibility, scalability and predictive performance with a probabilistic approach to measure uncertainty.

The challenge with bayesian neural networks is that we have to specify a meaningful prior distribution in advance and also the calculation of posterior distribution is intractable. A good prior distribution is difficult to get as the relationship between the weights of the network (Graves, 2011) and the output is non linear in nature while the calculation of the posterior requires doing an integral which is often intractable in nature (Kingma and Welling, 2013).

To avoid the above two difficulties, there has been considerable work done showing that as the width of a BNN was increased, the limiting distribution turns out to be a Gaussian process (Lee et al., 2019). However still the relationship of BNN with GP remains unclear as it fails to match the predictions made by GP. This could be due to the fact that there are a lot of kernels with different structured approximations and finding the one which best suits the task at hand is not straight forward.

In this paper, we perform variational inference with a new type of model architecture which we named as stochastic bayesian neural network. The update step is similar to the traditional backpropagation algorithms. In this method, a BNN is trained to produce a custom distribution with small KL-divergence with the true posterior. We do this by maximizing the Evidence Lower Bound (ELBO) by sampling based approximation. We specify stochastic process priors which are by their inherent nature rich in structured dependencies between function values. Using this method, we can model various structures including periodicity and smoothness (Sun et al., 2017). Thus stochastic bayesian

* Website of author - <https://abhinavsagar.github.io/>

neural networks combine the advantage of GP with the fact that posterior distribution becomes tractable. Our network beats the previous state of the art on regression datasets.

2 Related Work

Bayesian Neural Networks using Variational Inference approach has a rich history, first being applied by (Neal, 1993). The work was later extended by (Graves, 2011) using gaussian priors using covariance estimates. Later, (Kingma et al., 2015) proposed Variational Autoencoders for generative modelling using reparameterization technique. More work from (Flam-Shepherd et al., 2017) have used gaussian variational posteriors while (Huszár, 2017) have used normalizing flows for computing the posterior distribution. (Gal et al., 2017) have shown that dropouts can be approximated as an ensemble of neural networks in a bayesian setting. Neural networks with dropout were also interpreted as BNNs (Gal and Ghahramani, 2016) and (Gal et al., 2017). Local reparameterization trick (Kingma et al., 2015) proposed a new perspective by adding an additional parameter in the latent space after the encoder.

Stochastic variational inference uses a new technique by using update rules which resemble ordinary backpropagation (Graves, 2011) and (Blundell et al., 2015). The challenge with this approach is the fact that computing the posterior distributions is difficult as it is intractable in nature (Louizos and Welling, 2016) and (Shi et al., 2017). Some of the popular choices for priors are Radial Basis Kernel (RBF), gaussian, gaussian mixture distributions etc. Other priors, including log-uniform priors (Kingma et al., 2015) and horseshoe priors (Ghosh et al., 2018) have also been used successfully.

One common approach used in all of the previous works is that they used priors over the model parameters. The posterior distribution resulting is often intractable and also weight space distributions are difficult to characterize. In this paper, we used an alternative approach to automatically compute the prior using a well known theory known as stochastic process. The resulting neural networks which are still based on variational inference techniques are named as Stochastic Bayesian Neural Networks. Our method makes it possible to specify a range of priors and in particular stochastic process priors as has been done in gaussian process.

We summarize our main contributions as follows:

- An approach to take advantage of flexibility, scalability, predictive performance and a probabilistic approach to measure uncertainty of variational inference techniques on regression problems
- A theoretical analysis of our approach named Stochastic Bayesian Neural Network which uses an alternative lower bound which we call SELBO backed by stochastic process.
- Evaluation on the UCI dataset using test RMSE and log likelihood as the evaluation metrics shows we outperform previous state-of-the-art methods.

3 Background

3.1 Variational Inference

Bayesian neural networks are defined in terms of priors on weights and the likelihood of the observation. The goal in variational inference techniques is to maximize the ELBO with the goal of fitting an approximate posterior distribution (Blundell et al., 2015). Bayes by Backprop uses a fully factorized Gaussian approximation while computing the posterior distribution (Blundell et al., 2015). The gradients of ELBO can be computed by backpropagation using the local reparameterization trick by computing the gradients and using it for updates (Kingma and Welling, 2013).

Bayes theorem is used for finding the posterior, given the prior, evidence and likelihood as defined in Equation 1:

$$p(z | x) = \frac{p(x | z)p(z)}{p(x)} \quad (1)$$

However computation of the posterior distribution is infeasible due to the intractable integral in the likelihood term. This is where variational inference techniques come to rescue by converting

the equation to an optimization problem instead between the prior and posterior distributions. For measuring the difference between two probability distributions p and q , KL divergence is defined in Equation 2:

$$D_{KL}(q(x)\|p(x)) := E_{\sim q}[\Delta I] = \int (\Delta I)q(x)dx = \int q(x) \log \left(\frac{q(x)}{p(x)} \right) dx \quad (2)$$

The priors in variational inference techniques are chosen on the basis of computational convenience.

3.2 Variational Autoencoders

VAEs are a family of generative models which use an encoder-decoder architecture and have recently been used in a range of applications like generating images, generating music, recommender systems etc. The encoder converts the sampling distribution to a latent space in the form of mean and variance vectors, while the decoder reconstructs the original sample using both reconstruction error and the KL divergence between the prior and posterior distributions. Let posterior distribution in encoder be defined as $q(z|x)$, weights by θ and encoder as $q_\theta(z|x)$. Let the likelihood function in decoder be denoted as $p_\phi(x|z)$, and weights by ϕ .

The KL divergence between the approximate and the real posterior distributions is defined in Equation 3:

$$D_{KL}(q_\theta(z|x_i)\|p(z|x_i)) = - \int q_\theta(z|x_i) \log \left(\frac{p(z|x_i)}{q_\theta(z|x_i)} \right) dz \geq 0 \quad (3)$$

The above equation can be converted to an optimization problem as shown in Equation 4:

$$\log p(x_i) \geq -D_{KL}(q_\theta(z|x_i)\|p(z)) + E_{\sim q_\theta(z|x_i)} [\log p_\phi(x_i|z)] \quad (4)$$

The right hand side of the above equation is known as the Evidence Lower Bound (ELBO). The goal is to maximize the ELBO which maximizes the log probability. The first term in the above equation denotes the KL divergence between the true and approximate posterior distributions while the second term denotes the reconstruction error.

4 Proposed Method

Our method can be cast as two player zero sum game analogous to a generative adversarial network (GAN) (Goodfellow et al., 2014). Let the dataset be defined D , variational posterior $g()$, prior p , weight λ and sampling distribution s for random measurement points.

In this work, we used a sampling based approach. The network needs to match the prior distribution both near the training data and the test data where predictions are required. This is shown in Equation 5 and Equation 6 where X denotes the M samples independently drawn from c .

$$\text{Sample Points } \mathbf{X}^M \sim s \quad (5)$$

$$\mathbf{f}_i = g([\mathbf{X}^M]; \theta), i = 1 \dots k \quad (6)$$

The network is trained using stochastic gradient descent as shown in Equation 7:

$$\Delta = \frac{1}{k} \frac{1}{|D|} \sum_i \sum_{(x,y)} \nabla_\theta \log p(y|\mathbf{f}_i(x)) \quad (7)$$

Finally Adam optimizer is used for updating the posterior distribution using the prior distribution and the likelihood in every iteration until the distribution converges. This step is shown in Equation 8:

$$\phi \leftarrow \text{Optimizer}(\theta, \lambda\Delta) \quad (8)$$

Here λ is a regularization parameter which is tuned using bayesian optimization techniques. The optimal value of λ was found to be 0.24.

4.1 Stochastic Evidence Lower Bound (SELBO)

In our technique, we use a stochastic prior which can be any distribution including the well known Gaussian Process. Here we consider the neural network with stochastic weights and stochastic bias. We sample a function by sampling a random noise vector for some function. This sampling in turns helps in uncertainty quantification by maximizing the Stochastic Evidence Lower Bound (SELBO). The difference with the original ELBO is that in our case the distribution over the weights have been replaced by that over functions.

The KL term here represents the KL divergence between two stochastic processes instead of over two distributions. The computation of the KL-divergence between stochastic processes requires doing an integral which can be intractable in nature depending on the problem.

4.2 The Algorithm

Next, we present our algorithm used in this paper. In every iteration, we sample a mini batch of training data D and random points X from a distribution c . We forward the sample through a network $g(\phi)$ which defines the posterior distribution. The goal is to maximize the objective function defined which we name as Stochastic Evidence Lower Bound as shown in Equation 5:

$$\frac{1}{|\mathcal{D}|} \sum_{(\mathbf{x},y) \in \mathcal{D}_\theta} \mathbb{E}_{q_\theta} [\log p(y|f(\mathbf{x}))] - \lambda \text{KL} [q(\mathbf{f}^D) \| p(\mathbf{f}^D)] \quad (9)$$

Here λ is a regularization hyperparameter which needs to be tuned carefully to avoid overfitting.

Algorithm 1: Stochastic Bayesian Neural Networks (SBNN)

Input: Dataset D , variational posterior $g()$, prior p , weight λ
Sampling distribution s for random measurement points

```

while  $\theta$  not converged do
    Sample Points  $\mathbf{X}^M \sim s$ 
     $\mathbf{f}_i = g([\mathbf{X}^M]; \theta), i = 1 \dots k$ 
     $\Delta = \frac{1}{k} \frac{1}{|\mathcal{D}|} \sum_i \sum_{(x,y)} \nabla_\theta \log p(y|\mathbf{f}_i(x))$ 
     $\phi \leftarrow \text{Optimizer}(\theta, \lambda\Delta)$ 
end

```

4.3 Hyperparameters

The hyperparameters used in our model are specified in Table 1.

Table 1: Hyperparameters details

Parameter	Value
Batch Size	16
Optimizer	Adam
Learning Rate	0.0002

5 Results

Next we show our results in Table 1 and Table 2 respectively. We used 5 publicly available UCI datasets for regression and used two evaluation metrics - test RMSE and log likelihood for testing.

Table 2: Averaged test RMSE for the regression benchmarks

Dataset	(Blundell et al., 2015)	(Zhang et al., 2018)	Ours
Boston	3.171 ± 0.149	2.742 ± 0.125	2.424 ± 0.112
Concrete	5.678 ± 0.087	5.019 ± 0.127	5.003 ± 0.107
Energy	0.565 ± 0.018	0.485 ± 0.023	0.408 ± 0.019
Wine	0.643 ± 0.012	0.637 ± 0.011	0.653 ± 0.005

Table 3: Averaged log-likelihood for the regression benchmarks

Dataset	(Blundell et al., 2015)	(Zhang et al., 2018)	Ours
Boston	-2.602 ± 0.031	-2.446 ± 0.029	-2.296 ± 0.042
Concrete	-3.149 ± 0.018	-3.039 ± 0.025	-3.016 ± 0.015
Energy	-1.500 ± 0.006	-1.421 ± 0.005	-0.824 ± 0.017
Wine	-0.977 ± 0.017	-0.969 ± 0.014	-1.025 ± 0.014

6 Conclusions

In this paper, we investigated a new technique for training bayesian neural networks using stochastic processes. We proposed a new lower bound using variational inference techniques which we named as Stochastic Evidence Lower Bound. We train the neural network using gradient descent algorithms by sampling a mini batch of data in every iteration. Using test RMSE and log likelihood as the evaluation metrics, our work outperforms previous state of the art on 5 publicly available UCI datasets on regression benchmarks. This approach allows estimating uncertainties while also being scalable to larger datasets.

Acknowledgments

We would like to thank Nvidia for providing the GPUs.

References

- M. Arjovsky and L. Bottou. Towards principled methods for training generative adversarial networks. *arXiv preprint arXiv:1701.04862*, 2017.
- A. Asuncion and D. Newman. Uci machine learning repository, 2007.
- C. Blundell, J. Cornebise, K. Kavukcuoglu, and D. Wierstra. Weight uncertainty in neural networks. *arXiv preprint arXiv:1505.05424*, 2015.
- S. Depeweg, J. M. Hernández-Lobato, F. Doshi-Velez, and S. Udluft. Learning and policy search in stochastic dynamical systems with bayesian neural networks. *arXiv preprint arXiv:1605.07127*, 2016.
- D. Flam-Shepherd, J. Requeima, and D. Duvenaud. Mapping gaussian process priors to bayesian neural networks. In *NIPS Bayesian deep learning workshop*, 2017.
- Y. Gal and Z. Ghahramani. Dropout as a bayesian approximation: Representing model uncertainty in deep learning. In *international conference on machine learning*, pages 1050–1059, 2016.
- Y. Gal, R. Islam, and Z. Ghahramani. Deep bayesian active learning with image data. *arXiv preprint arXiv:1703.02910*, 2017.
- S. Ghosh, J. Yao, and F. Doshi-Velez. Structured variational learning of bayesian neural networks with horseshoe priors. *arXiv preprint arXiv:1806.05975*, 2018.
- I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio. Generative adversarial nets. In *Advances in neural information processing systems*, pages 2672–2680, 2014.

- A. Graves. Practical variational inference for neural networks. In *Advances in neural information processing systems*, pages 2348–2356, 2011.
- J. M. Hernández-Lobato and R. Adams. Probabilistic backpropagation for scalable learning of bayesian neural networks. In *International Conference on Machine Learning*, pages 1861–1869, 2015.
- F. Huszár. Variational inference using implicit distributions. *arXiv preprint arXiv:1702.08235*, 2017.
- M. E. Khan, D. Nielsen, V. Tangkaratt, W. Lin, Y. Gal, and A. Srivastava. Fast and scalable bayesian deep learning by weight-perturbation in adam. *arXiv preprint arXiv:1806.04854*, 2018.
- D. P. Kingma and M. Welling. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*, 2013.
- D. P. Kingma, T. Salimans, and M. Welling. Variational dropout and the local reparameterization trick. In *Advances in neural information processing systems*, pages 2575–2583, 2015.
- J. Lee, L. Xiao, S. Schoenholz, Y. Bahri, R. Novak, J. Sohl-Dickstein, and J. Pennington. Wide neural networks of any depth evolve as linear models under gradient descent. In *Advances in neural information processing systems*, pages 8572–8583, 2019.
- Y. Li and Y. Gal. Dropout inference in bayesian neural networks with alpha-divergences. *arXiv preprint arXiv:1703.02914*, 2017.
- C. Louizos and M. Welling. Structured and efficient variational deep learning with matrix gaussian posteriors. In *International Conference on Machine Learning*, pages 1708–1716, 2016.
- C. Ma, Y. Li, and J. M. Hernández-Lobato. Variational implicit processes. In *International Conference on Machine Learning*, pages 4222–4233, 2019.
- W. J. Maddox, P. Izmailov, T. Garipov, D. P. Vetrov, and A. G. Wilson. A simple baseline for bayesian uncertainty in deep learning. In *Advances in Neural Information Processing Systems*, pages 13153–13164, 2019.
- A. G. d. G. Matthews, J. Hensman, R. Turner, and Z. Ghahramani. On sparse variational methods and the kullback-leibler divergence between stochastic processes. *Journal of Machine Learning Research*, 51:231–239, 2016.
- J. Mukhoti, P. Stenetorp, and Y. Gal. On the importance of strong baselines in bayesian deep learning. *arXiv preprint arXiv:1811.09385*, 2018.
- R. M. Neal. Bayesian learning via stochastic dynamics. In *Advances in neural information processing systems*, pages 475–482, 1993.
- R. M. Neal. *Bayesian learning for neural networks*, volume 118. Springer Science & Business Media, 2012.
- G. Roeder, Y. Wu, and D. K. Duvenaud. Sticking the landing: Simple, lower-variance gradient estimators for variational inference. In *Advances in Neural Information Processing Systems*, pages 6925–6934, 2017.
- A. Sagar. Bayesian neural network via stochastic gradient descent. *arXiv preprint arXiv:2006.08453*, 2020.
- H. Salimbeni and M. Deisenroth. Doubly stochastic variational inference for deep gaussian processes. In *Advances in Neural Information Processing Systems*, pages 4588–4599, 2017.
- J. Shi, S. Sun, and J. Zhu. Kernel implicit variational inference. *arXiv preprint arXiv:1705.10119*, 2017.
- S. Sun, C. Chen, and L. Carin. Learning structured weight uncertainty in bayesian neural networks. In *Artificial Intelligence and Statistics*, pages 1283–1292, 2017.

S. Sun, G. Zhang, J. Shi, and R. Grosse. Functional variational bayesian neural networks. *arXiv preprint arXiv:1903.05779*, 2019.

H. Wang and D.-Y. Yeung. Towards bayesian deep learning: A framework and some existing methods. *IEEE Transactions on Knowledge and Data Engineering*, 28(12):3395–3408, 2016.

G. Zhang, S. Sun, D. Duvenaud, and R. Grosse. Noisy natural gradient as variational inference. In *International Conference on Machine Learning*, pages 5852–5861, 2018.